

Chapter 4: Logistic Regression and Classification in R

Syntax Notes, Exercises, and Worked Solutions

STAT452 – Applied Statistics for Engineers and Scientists II

Contents

1	What this worksheet covers	1
2	R syntax	1
3	Exercises	1
4	Worked solutions	2

1 What this worksheet covers

For a binary response $Y \in \{0, 1\}$, linear regression can produce values outside $[0, 1]$. Logistic regression models the log-odds: $\log\{\pi(x)/(1 - \pi(x))\} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$, where $\pi(x) = P(Y = 1 \mid x)$. R fits the coefficients by maximum likelihood with `glm(..., family=binomial)`. Exponentiating a coefficient gives an odds ratio.

2 R syntax

```
fit <- glm(transmission ~ wt + hp, data=cars, family=binomial)
prob <- predict(fit, type="response")
pred <- ifelse(prob >= 0.5, "manual", "automatic")
exp(coef(fit))
anova(null_fit, fit, test="Chisq")
```

The 0.5 threshold is a decision rule, not a universal law. Report a confusion matrix and metrics such as accuracy, sensitivity, and specificity, preferably on held-out data.

3 Exercises

Exercise 1: Compare with linear regression. Regress binary `am` on `wt` and `hp`. Check whether fitted values lie in $[0, 1]$.

Exercise 2: Fit and interpret the logit model. Fit logistic regression. Compute odds ratios and explain coefficient signs while holding the other predictor fixed.

Exercise 3: Classify observations. Convert probabilities to classes at threshold 0.5. Produce a confusion matrix and calculate accuracy.

Exercise 4: Threshold trade-offs. Compare thresholds 0.3 and 0.5. Calculate sensitivity and specificity.

Exercise 5: Validate predictions. Split into training and test observations. Fit on training data and evaluate on test data. Explain optimistic training accuracy.

Exercise 6: Compare models. Use a likelihood-ratio test and AIC to compare intercept-only and two-predictor models.

4 Worked solutions

Solution 1. Linear regression is not guaranteed to respect probability bounds; inspect `range(predict(linear_binary))`. This motivates the logistic link. □

Solution 2. Use `exp(coef(fit))`. An odds ratio above one raises the odds of the modeled success class; below one lowers the odds, conditional on other predictors. □

Solution 3. A confusion matrix separates correct and incorrect classifications. Accuracy is $(TP + TN)/n$, but it can hide unequal costs or class imbalance. □

Solution 4. Lowering the threshold generally increases sensitivity and decreases specificity. Choose it based on the cost of false negatives versus false positives. □

Solution 5. Test observations were not used during estimation, so test performance better reflects generalization. With 32 cars, treat this split as practice rather than a definitive performance claim. □

Solution 6. The likelihood-ratio test compares nested models using change in deviance. AIC balances likelihood and parameter count; lower AIC is preferred among models compared. □

Checklist

`glm -> predict(type=response) -> threshold -> confusion matrix -> held-out validation.`